PitchReady

Financial Estimation Report

Full infrastructure cost model across MVP, Growth, and Scale stages · March 2026

1. Overview

This report covers every financial dependency in the PitchReady stack — from day-one infrastructure to per-user AI costs — across three growth stages. Use this to understand your burn rate before the first line of code is written, and to model the point at which revenue covers operating costs.

All figures are in USD. AI cost estimates are based on Anthropic claude-sonnet-4-5 pricing as of March 2026. Estimates assume moderate, realistic usage — not worst-case.

2. Service-by-Service Breakdown

2.1 Vercel — Frontend Hosting

Plan	Cost	Bandwidth	Suitable For
Hobby (Free)	$0 / mo	100 GB	MVP & early testing
Pro	$20 / mo	1 TB	Growth stage, custom domains
Recommended start	**$0 / mo**	**—**	**Free until 100+ active users**

Note: Hobby plan does not support SLA. Upgrade to Pro once you have paying users.

2.2 DigitalOcean Droplet — Backend + Worker + Nginx

Hosts FastAPI, LangGraph worker, and Nginx reverse proxy inside Docker. This is your primary monthly fixed cost.

Droplet Size	vCPU	RAM	Cost/mo	Stage
Basic (s-1vcpu-2gb)	1	2 GB	$12 / mo	MVP

Droplet Size	vCPU	RAM	Cost/mo	Stage
Regular (s-2vcpu-4gb)	2	4 GB	$24 / mo	Growth
Regular (s-4vcpu-8gb)	4	8 GB	$48 / mo	Scale
Recommended start	**1**	**2 GB**	**$12 / mo**	**MVP**

Upgrade trigger: CPU consistently above 70% or worker queue latency exceeds 60 seconds.

2.3 DigitalOcean Spaces — File Storage (PDF / PPTX)

Stores all uploaded pitch decks. Users upload directly via pre-signed URLs — no file bytes pass through your server.

Component	Included	Overage	Est. Monthly
Base plan	250 GB storage	—	$5.00
Outbound transfer	1 TB / mo	$0.01 / GB	~$0.00 (MVP)
Extra storage	—	$0.02 / GB	~$0.00 (MVP)
Total (MVP)	**—**	**—**	**$5.00 / mo**

Avg pitch deck: 5–15 MB. 1,000 uploads ≈ 10 GB storage — well within the base plan.

2.4 Neon — Serverless Postgres + Auth

Handles all persistent data (users, sessions, analysis runs, reports) and authentication. The free tier is generous enough for MVP.

Plan	Storage	Compute Hours	Cost
Free	0.5 GB	191 hrs / mo	$0 / mo
Launch	10 GB	300 hrs / mo	$19 / mo
Scale	50 GB	750 hrs / mo	$69 / mo
Recommended start	**Free tier**	**—**	**$0 / mo**

Upgrade trigger: DB storage exceeds 400 MB or compute hours run out mid-month (visible in Neon dashboard).

2.5 Upstash Redis — Job Queue

Manages the async analysis job queue between FastAPI and the LangGraph worker. Serverless — you pay per command, not per instance.

Tier	Commands / Day	Price	Est. Monthly
Free	10,000 / day	$0	$0.00
Pay-as-you-go	Unlimited	$0.2 / 100K cmds	~$0.50–$2.00
Recommended start	**Free tier**	**—**	**$0 / mo**

Each analysis run uses ~20–30 Redis commands. Free tier covers ~300–500 runs/day — more than sufficient for MVP.

2.6 Anthropic Claude API — AI Engine

The primary variable cost in your stack. Every analysis run makes 4 sequential API calls. Costs scale linearly with usage.

Agent	Input Tokens	Output Tokens	Cost/Run	Notes
Agent 1 — Extractor	~3,000	~800	~$0.021	Parses deck to JSON
Agent 2 — Assumption Mapper	~2,000	~600	~$0.015	Ranks unproven claims
Agent 3 — Domain Interrogator	~3,500	~1,000	~$0.026	Vertical KB injected
Agent 4 — Report Synthesizer	~4,000	~1,500	~$0.035	Final report output
Total per run	**~12,500**	**~3,900**	**~$0.097**	**≈ $0.10 / run**

If Agent 3 vertical knowledge base grows large (8,000–15,000 tokens), cost per run rises to $0.18–$0.28. Still manageable.

API pricing reference (claude-sonnet-4-5):

- Input tokens: $3.00 per million tokens
- Output tokens: $15.00 per million tokens

2.7 PostHog — User Analytics

Plan	Events / Mo	Features	Cost
Free	1 million	Analytics, replays, flags	$0 / mo
Paid	Unlimited	+ priority support	$0.00031 / event
Recommended start	**Free tier**	**—**	**$0 / mo**

1 million events/month is roughly 10,000 active users generating 100 events each. You won't hit this for a long time.

2.8 Domain Name

Item	Cost	Notes
.com domain (e.g. Namecheap)	~$12–15 / yr	$1.00–1.25 / mo amortized
SSL Certificate	$0	Certbot + Let's Encrypt, free
Total	**~$13 / yr**	**~$1.08 / mo**

3. Stage-by-Stage Monthly Cost Model

Three stages are modelled: MVP (0–100 users), Growth (100–1,000 users), and Scale (1,000–5,000 users). User count refers to registered users, not concurrent sessions.

3.1 MVP Stage — 0 to 100 Users

Assumes: 100 registered users, 30% active monthly, avg 3 runs per active user = ~90 runs/month.

Service	Monthly Cost	Notes
Vercel (Hobby)	$0.00	Free tier
DO Droplet (1 vCPU / 2 GB)	$12.00	Backend + worker
DO Spaces	$5.00	Base plan
Neon (Free)	$0.00	Free tier
Upstash Redis (Free)	$0.00	Free tier
Claude API (~90 runs × $0.10)	$9.00	Variable
PostHog (Free)	$0.00	Free tier
Domain (amortized)	$1.08	~$13/yr
TOTAL	**$27.08 / mo**	**~$325 / year**

3.2 Growth Stage — 100 to 1,000 Users

Assumes: 1,000 registered users, 35% active monthly, avg 5 runs per active user = ~1,750 runs/month.

Service	Monthly Cost	Notes
Vercel (Pro)	$20.00	Custom domain + SLA
DO Droplet (2 vCPU / 4 GB)	$24.00	Handles higher worker load
DO Spaces	$5.00	~8 GB used, within base plan
Neon (Launch)	$19.00	Storage + compute upgrade
Upstash Redis (Pay-as-you-go)	$1.50	~700K commands/mo
Claude API (~1,750 runs × $0.10)	$175.00	Primary variable cost
PostHog (Free)	$0.00	Still within 1M events
Domain (amortized)	$1.08	—

Service	Monthly Cost	Notes
TOTAL	**$245.58 / mo**	**~$2,947 / year**

3.3 Scale Stage — 1,000 to 5,000 Users

Assumes: 5,000 registered users, 40% active monthly, avg 6 runs per active user = ~12,000 runs/month.

Service	Monthly Cost	Notes
Vercel (Pro)	$20.00	—
DO Droplet (4 vCPU / 8 GB)	$48.00	May need 2 droplets for workers
DO Spaces	$7.00	~100 GB used
Neon (Scale)	$69.00	High compute hours
Upstash Redis	$5.00	~2.5M commands/mo
Claude API (~12,000 runs × $0.10)	$1,200.00	Dominant cost at scale
PostHog (Paid)	$15.00	~50M events/mo
Domain (amortized)	$1.08	—
TOTAL	**$1,365.08 / mo**	**~$16,381 / year**

4. Revenue vs. Cost — Breakeven Analysis

Using a subscription model at $29/month (the recommended pricing tier). Breakeven is the number of paying subscribers needed to cover total monthly infrastructure costs.

Stage	Monthly Cost	Paying Users to Break Even	At $29/mo
MVP	$27.08	1 paying user	1 × $29 = $29
Growth	$245.58	9 paying users	9 × $29 = $261
Scale	$1,365.08	48 paying users	48 × $29 = $1,392

The unit economics are very healthy. At scale with 5,000 users, even a 1% conversion rate (50 paying users) covers the entire infrastructure bill. A 5% conversion rate generates $7,250/month against $1,365 in costs — a 5.3× return on infrastructure spend.

5. Free Tier Cost — First 3 Uses

Each new user gets 3 free analysis runs before hitting a paywall. Here is what that costs you:

Item	Calculation	Cost
Claude API (3 runs)	3 × $0.10	$0.30
DO Spaces storage	~30 MB per user	~$0.001
Compute (Droplet)	Amortized per run	~$0.005
Neon DB writes	Negligible	~$0.00
Cost to acquire one trial user	**—**	**~$0.31**

Your customer acquisition cost from an infrastructure perspective is thirty-one cents. This is exceptionally low — most SaaS products spend $5–50 in infrastructure costs per trial user. The 3-run free tier is essentially free to offer.

6. Cost Optimization Tips

Short-term (MVP)

- Cache Agent 1 (Extractor) output — if the same deck is re-analyzed, skip re-extraction and save ~$0.021 per run
- Set a max token cap on Agent 3's vertical knowledge base injection to control the biggest cost variable
- Log tokens per run from day one — you want data before costs surprise you

Medium-term (Growth)

- Batch low-priority re-analysis jobs during off-peak hours using Upstash Queue — Anthropic doesn't charge differently for timing, but your Droplet stays cooler
- Consider prompt caching for Agent 3 system prompts — Anthropic's prompt caching feature reduces repeated long-context costs by up to 90%
- Move completed report files out of hot Spaces storage to cheaper archival after 90 days

Long-term (Scale)

- Negotiate Anthropic volume pricing — available above $5,000/month spend
- Run Agent 1 (Extractor) locally with a smaller open-source model (Mistral 7B) to cut ~20% off per-run cost
- Consider splitting the worker Droplet from the API Droplet for independent scaling

---

7. Summary

Stage	Users	Monthly Cost	Break-even Subscribers
MVP	0 – 100	$27 / mo	1
Growth	100 – 1,000	$246 / mo	9
Scale	1,000 – 5,000	$1,365 / mo	48

PitchReady has exceptionally lean unit economics. The dominant cost at every stage is the Claude API — which scales proportionally with usage and is directly offset by the value delivered per run. Infrastructure overhead is low, break-even is reachable with single-digit paying users, and the free tier costs less than a dollar per user to offer.